

Contents

Chinese AI Coding Models, Tools & Platforms: Comprehensive Guide (March 2026)	1
Table of Contents	2
1. Models by Company	2
DeepSeek	2
Alibaba / Qwen	4
Moonshot AI / Kimi	5
Zhipu AI (Z.ai)	6
MiniMax	7
01.AI (Yi)	8
Baidu (ERNIE)	8
Ant Group (CodeFuse / Ling)	9
SenseTime (SenseNova)	10
Baichuan	10
Shanghai AI Lab (InternLM)	11
2. IDE & Coding Tools	11
Tongyi Lingma (Alibaba)	11
Qwen Code CLI (Alibaba)	11
Kimi Code CLI (Moonshot)	12
CodeGeeX (Zhipu AI)	12
Trae IDE (ByteDance)	12
Baidu Comate	13
Tencent CodeBuddy	13
Fitten Code	14
DevChat (Merico)	14
3. Benchmarks Comparison Table	14
Coding Benchmarks: Chinese vs. Western Models (March 2026)	14
4. Unique Strengths of Chinese Models	15
4.1 Training Efficiency & Cost	15
4.2 Open-Source Leadership	15
4.3 MoE Architecture Innovation	15
4.4 Multilingual & Chinese-Language Superiority	16
4.5 Agentic Coding Innovation	16
4.6 Vertical Integration	16
5. Access & Availability Outside China	16
Freely Available Globally	16
Restricted or China-Focused	16
Third-Party Inference Providers	17
Integration with Western Developer Tools	17
6. Chinese Coding Benchmarks & Communities	17
Benchmarks	17
Open-Source Communities & Platforms	17
7. Regulatory Considerations	18
Data Sovereignty	18
US Export Controls Impact	18
Open-Source as Strategic Policy	18
Implications for Western Developers	18
Summary	18

Chinese AI Coding Models, Tools & Platforms: Comprehensive Guide (March 2026)

The Chinese AI coding ecosystem has undergone a dramatic transformation since 2024. What was once a 30+ percentage-point gap on HumanEval has collapsed to near parity – and in several categories, Chinese

models now lead. This document covers every major model, tool, platform, and ecosystem development as of March 2026.

Table of Contents

- 1. Models by Company
 - DeepSeek
 - Alibaba / Qwen
 - Moonshot AI / Kimi
 - Zhipu AI (Z.ai)
 - MiniMax
 - 01.AI (Yi)
 - Baidu (ERNIE)
 - Ant Group (CodeFuse / Ling)
 - SenseTime (SenseNova)
 - Baichuan
 - Shanghai AI Lab (InternLM)
- 2. IDE & Coding Tools
 - Tongyi Lingma (Alibaba)
 - Qwen Code CLI (Alibaba)
 - Kimi Code CLI (Moonshot)
 - CodeGeeX (Zhipu AI)
 - Trae IDE (ByteDance)
 - Baidu Comate
 - Tencent CodeBuddy
 - Fitten Code
 - DevChat (Merico)
- 3. Benchmarks Comparison Table
- 4. Unique Strengths of Chinese Models
- 5. Access & Availability Outside China
- 6. Chinese Coding Benchmarks & Communities
- 7. Regulatory Considerations

1. Models by Company

DeepSeek

DeepSeek, founded in 2023 and funded by the quantitative trading firm High-Flyer, has become the most prominent Chinese AI lab in global developer consciousness. Their models are fully open-source under the MIT License.

DeepSeek-Coder-V2 (June 2024)

Attribute	Details
Creator	DeepSeek
Release	June 2024
License	DeepSeek License (commercial use allowed)
Architecture	MoE (Mixture-of-Experts)
Total Parameters	236B (large) / 16B (lite)
Active Parameters	21B (large) / 2.4B (lite)
Context Window	128K tokens
Languages Supported	338 programming languages
Training Data	6T additional tokens on top of DeepSeek-V2
HumanEval	~85% (Pass@1)
Availability	HuggingFace, ModelScope, API

DeepSeek-Coder-V2 was the first open-source model to match GPT-4 Turbo on code-specific tasks. It was further pre-trained from an intermediate checkpoint of DeepSeek-V2, substantially enhancing coding and mathematical reasoning while maintaining general language capabilities.

DeepSeek-V3 (December 2024)

Attribute	Details
Creator	DeepSeek
Release	December 2024
License	MIT License
Architecture	MoE with Multi-head Latent Attention (MLA)
Total Parameters	671B
Active Parameters	37B per token
Context Window	128K tokens
Training Cost	~\$5.6M (2.788M H800 GPU hours)
HumanEval	85.7%
MBPP	75.4%
SWE-bench Verified	47.8%

Key Innovations: - **FP8 Mixed Precision Training:** First model to validate FP8 training at extreme scale, reducing both GPU memory and compute. - **Multi-Head Latent Attention (MLA):** Reduces KV-cache memory by 93.3%, dramatically cutting inference costs. - **Auxiliary-Loss-Free Load Balancing:** Eliminates auxiliary losses for MoE routing, using bias terms instead – a novel approach that improves both training stability and quality. - **Training Cost:** The reported \$5.6M figure covers only GPU compute for pre-training. Total R&D costs (including prior research, failed experiments, and team salaries) are estimated at \$500M-\$1.3B. Regardless, the per-token training efficiency is extraordinary.

DeepSeek-R1 (January 2025)

Attribute	Details
Creator	DeepSeek
Release	January 20, 2025
License	MIT License
Architecture	Same 671B MoE as V3, with RL-enhanced reasoning
Context Window	64K tokens
API Pricing	\$0.55/M input, \$2.19/M output tokens
HumanEval-Mul	82.6%
SWE-bench Verified	50.8%
MATH-500	90.2%
Codeforces Rating	2,029 (Candidate Master)

R1 was trained with heavy reinforcement learning to develop Chain-of-Thought reasoning, reflection, and multi-step verification. It matches or exceeds OpenAI’s o1 on most mathematical and reasoning benchmarks.

DeepSeek-V3.1 (August 2025)

Attribute	Details
Release	August 2025
License	MIT License
Key Improvements	Optimized for domestic AI accelerators; hybrid thinking/non-thinking modes
Aider Benchmark	71.6% pass rate (surpassing Claude Opus)

Attribute	Details
Cost	68x cheaper than Claude Opus

V3.1 marked DeepSeek’s deliberate pivot toward hardware independence and supply chain resilience. Its Deep Thinking Mode achieves 90-95% of R1’s reasoning performance.

DeepSeek-V3.2 (December 2025)

Attribute	Details
Release	December 2025
License	MIT License
Key Innovation	DeepSeek Sparse Attention (DSA) – O(n) complexity
AIME	96.0% (V3.2-Speciale)
Competition Results	Gold-level at IMO, CMO, ICPC World Finals, IOI 2025

V3.2-Speciale outperforms GPT-5 on several reasoning benchmarks. It is the first model to integrate thinking directly into tool-use, supporting both thinking and non-thinking modes during tool calls.

Alibaba / Qwen

Alibaba’s Qwen (Tongyi Qianwen) team has built the most prolific open-source model family in the world. By mid-2025, Qwen became the model with the most derivatives on HuggingFace, surpassing Meta’s Llama with over 113,000 derivative models.

Qwen2.5-Coder (November 2024)

Attribute	Details
Creator	Alibaba / Qwen Team
Release	November 2024
License	Apache 2.0 (most sizes), Qwen License (larger)
Architecture	Dense Transformer
Sizes	0.5B, 1.5B, 3B, 7B, 14B, 32B
Context Window	128K tokens
Training Data	5.5T tokens, 92+ programming languages
Variants	Base and Instruct for each size
Availability	HuggingFace, ModelScope, Ollama

HumanEval Pass@1 scores by size (Instruct variants):

Size	HumanEval	Notes
0.5B	~61%	Smallest competitive code model
1.5B	70.7%	
3B	78.7%	
7B	79.7%	Surpasses CodeStral-22B
14B	80.5%	Surpasses DS-Coder-33B-Instruct
32B	92.7% (base) / 88.4% (instruct)	SOTA among open-source

Additional 32B-Instruct scores: - MBPP: 84.0% - LiveCodeBench: 51.2% - BigCodeBench, EvalPlus, Cross-CodeEval: SOTA among open-source

The Qwen2.5-Coder series demonstrates remarkable parameter efficiency – the 7B model outperforms models 3-4x its size.

Qwen3-Coder (Mid-2025)

Attribute	Details
Creator	Alibaba / Qwen Team
Release	Mid-2025
Architecture	MoE
Total Parameters	480B
Active Parameters	35B per token
SWE-bench Verified	~69.6%
Key Feature	Performs on par with Claude Sonnet 4 on agentic coding

Accompanied by the Qwen Code CLI tool, Qwen3-Coder set new benchmarks among open models for agentic coding, browser use, and tool use.

Qwen3-Coder-Next (February 2026)

Attribute	Details
Creator	Alibaba / Qwen Team
Release	February 2026
Architecture	Ultra-sparse MoE
Total Parameters	80B
Active Parameters	3B per token
SWE-bench Verified	70.6%
SWE-bench Pro	44.3%
SWE-bench Multilingual	62.8%
Throughput	Up to 10x higher than dense models of similar size

Designed specifically for coding agents and local development. Trained on ~800K verifiable tasks with executable environments, using large-scale RL. Outperforms Claude Opus 4.5 on secure code generation benchmarks.

Moonshot AI / Kimi

Moonshot AI, founded by former Tsinghua and Google researchers, has rapidly risen to become one of China’s most capable AI labs, particularly in agentic coding.

Kimi K2 (July 2025)

Attribute	Details
Creator	Moonshot AI
Release	July 2025
License	Modified MIT License
Architecture	MoE with MLA
Total Parameters	1 trillion
Active Parameters	32B
Layers	61 (including 1 dense layer)

Attribute	Details
Experts	384 total, 8 selected + 1 shared per token
Context Window	128K tokens
SWE-bench Verified	65.8%
GSM8K	94.4% (surpassing GPT-4.1 and Claude Opus 4)

Kimi K2.5 (January 2026)

Attribute	Details
Creator	Moonshot AI
Release	January 27, 2026
License	Open-weight
Architecture	1T MoE + MoonViT-3D (400M vision encoder)
Context Window	256K tokens
SWE-bench Verified	76.8%
AIME 2025	96.1%
Humanity’s Last Exam	50.2% (at 76% lower cost than Claude Opus 4.5)
API Pricing	\$0.60/M input tokens
Availability	HuggingFace, NVIDIA NIM, Kimi API

Key Innovation – Agent Swarm Technology: Coordinates up to 100 specialized AI agents working simultaneously, cutting execution time by 4.5x. Uses Parallel Agent Reinforcement Learning (PARL), a novel RL technique that addresses training instability, credit assignment, and “serial collapse” in multi-agent settings.

Visual-to-Code: As a native multimodal model, K2.5 turns text, images, and video inputs into functional front-end code with high fidelity. Powers complete visual-to-code workflows including complex page layouts, full websites with animations, and interactive UI.

Industry Impact: Cursor confirmed that its Composer 2 was built partially on Kimi technology (~25% of compute from Kimi base model).

Zhipu AI (Z.ai)

Zhipu AI, spun out of Tsinghua University’s KEG Lab, rebranded as Z.ai in late 2025. Their GLM series has become a serious contender in coding.

GLM-4.7 (December 2025)

Attribute	Details
Creator	Zhipu AI / Z.ai
Release	December 22, 2025
License	Open-weight
Architecture	MoE
Total Parameters	~355B
Active Parameters	32B per forward pass
Context Window	200K tokens (128K max output)
Inference Speed	55 tokens/sec
API Pricing	~\$3/month

Benchmark Scores:

Benchmark	Score
HumanEval	94.2%
SWE-bench Verified	73.8%
LiveCodeBench v6	84.9%
AIME 2025	95.7%
MATH	92%
GSM8k	98%
MMLU	88%

Advanced Features: - **Preserved Thinking:** Maintains reasoning chains across multiple conversation turns. - **Interleaved Thinking:** Pauses to “reason aloud” before executing tool calls.

GLM-4.7 is arguably the most well-rounded open-source coding model available as of early 2026, matching premium competitors while costing a fraction.

Industry Impact: Windsurf confirmed that its core model (SWE-1.5) originated from Z.ai / Zhipu AI technology.

GLM-5 (February 2026)

Attribute	Details
Release	February 2026
Pricing	~30% higher than GLM-4.7
Improvements	Enhanced reasoning and code generation

MiniMax

MiniMax, founded in 2021 by former SenseTime VP Yan Junjie, has made an aggressive push into coding with its M-series models.

MiniMax-M2 (Q4 2025)

Attribute	Details
Release	Q4 2025
License	Open-source
Key Focus	Coding, tool use, deep search
Integration	Claude Code, Cursor, Cline, Kilo Code, Droid

MiniMax-M2.5 (February 2026)

Attribute	Details
Creator	MiniMax
Release	February 12, 2026
License	Open-source
Architecture	MoE with Lightning Attention, 8 experts per token
Total Parameters	230B
Active Parameters	10B
SWE-bench Verified	80.2% (#1 among open-source, ahead of Claude Opus 4.6)
Multi-SWE-Bench	51.3% (#1 overall, ahead of Opus 4.6’s 50.3%)
BrowseComp	76.3%

Attribute	Details
HumanEval	89.6%
Training Languages	10+ (Go, C, C++, TypeScript, Rust, Kotlin, Python, Java, JS, PHP, Lua, Dart, Ruby)
Training Environments	200,000+ real-world RL environments

M2.5 plans features like an architect – actively decomposing tasks, planning structure and UI before writing code. It’s 37% faster than M2.1 on agentic tasks. At 10B active parameters, it delivers frontier performance at roughly 1/20th the cost.

01.AI (Yi)

01.AI was founded by Dr. Kai-Fu Lee (former president of Google China, Microsoft Research, and Apple).

Yi-Coder (September 2024)

Attribute	Details
Creator	01.AI
Release	September 2024
License	Apache 2.0
Architecture	Dense Transformer (based on Yi-9B)
Sizes	1.5B, 9B
Context Window	128K tokens
Training Data	2.4T tokens, 52 programming languages
Availability	HuggingFace, Ollama

Yi-Coder-9B-Chat Benchmarks:

Benchmark	Score
HumanEval	85.4%
MBPP	73.8%
LiveCodeBench	23.4% (first <10B model to exceed 20%)
CRUXEval-O	>50% (first open-source to pass this threshold)
PAL Math	70.3% (surpasses DeepSeek-Coder-33B)

Yi-Coder proves that sub-10B models can deliver serious coding performance. However, the company’s focus has shifted more toward general-purpose models and consumer products (Yi Lightning, Yi Chat) in 2025-2026.

Baidu (ERNIE)

Baidu, often called “China’s Google,” has a deep AI research tradition through its PaddlePaddle framework and ERNIE model series.

ERNIE 4.5 (March 2025)

Attribute	Details
Creator	Baidu
Release	March 2025
License	Apache 2.0

Attribute	Details
Architecture	MoE
Sizes	10 variants from 0.3B dense to 424B MoE
Active Parameters	3B to 47B depending on variant
Context Window	128K tokens
HumanEval+	>90%
MBPP	>80%

Caveats: Independent assessments show ERNIE 4.5 performs significantly worse than GPT-4.5 and DeepSeek on LiveCodeBench and more challenging coding benchmarks. Its strengths are more in general knowledge, content creation, and Chinese-language tasks.

ERNIE 5.0 (November 2025) Announced at Baidu World 2025 with improved reasoning capabilities. Baidu reports that 45% of new code within Baidu is now AI-generated.

ERNIE X1.1 (Reasoning Model) A dedicated reasoning model with upgrades in mathematical computation, code generation, and tool use. Supports 128K context window.

Ant Group (CodeFuse / Ling)

Ant Group, Alibaba’s fintech arm (Alipay parent), operates a distinct AI research organization.

CodeFuse (2023-2025)

Attribute	Details
Creator	Ant Group
Release	September 2023 (initial), ongoing
License	Open-source
Focus	Full SDLC: design, requirements, coding, testing, deployment, ops
Notable Research	Code Graph Model (CGM) – accepted at NeurIPS 2025

CGM is a graph-based LLM framework for repository-level software engineering, using a chain of four atomic nodes (Rewriter, Retriever, Reranker, Reader) to navigate complex codebases.

Ling-1T (October 2025)

Attribute	Details
Creator	Ant Group
Release	October 2025
License	Open-source
Architecture	MoE with Evolutionary Chain-of-Thought (Evo-CoT)
Total Parameters	1 trillion
Active Parameters	~50B per token
Training Data	20T+ reasoning-dense tokens
Context Window	128K tokens
LiveCodeBench	Leading among 1T models (surpasses Kimi K2 at 53.7%, GPT-5 at 44.7%)

Attribute	Details
MMLU	91.76% (surpasses GPT-5, Claude 4.5 Sonnet, DeepSeek V3.1)
AIME 2025	70.42%
Availability	HuggingFace (inclusionAI/Ling-1T)

Ling-1T is the first flagship non-thinking model in the Ling 2.0 series, part of a family comprising Ling (non-thinking), Ring (thinking), and Ming (multimodal).

SenseTime (SenseNova)

SenseTime, originally a computer vision company, has expanded into general AI.

SenseNova V6 (April 2025)

Attribute	Details
Creator	SenseTime
Release	April 2025
Architecture	Hybrid MoE
Key Strengths	Multimodal reasoning, long CoT, mathematics, coding
Benchmarks	Multimodal reasoning ranked #1 in China (vs GPT-o1)
Cost	Lowest training and reasoning costs in Chinese industry
Notable	First Chinese model supporting 10-minute mid-to-long video analysis

SenseNova is stronger in multimodal and vision tasks than in pure coding. SenseTime reported record revenue of over RMB 5 billion in 2025, with EBITDA turning positive in H2.

Baichuan

Baichuan Intelligence, founded by former Sogou CEO Wang Xiaochuan, has pivoted primarily toward medical AI.

Baichuan-M3-235B (2025)

Attribute	Details
Creator	Baichuan Intelligence
Architecture	MoE (235B total)
Primary Focus	Medical AI, with general coding as secondary
Coding Benchmarks	Competitive at its scale on HumanEval/MBPP, but behind dedicated coding models
Availability	HuggingFace

Baichuan models support coding tasks but are not specialized for them. The company's primary differentiation is in medical AI (Baichuan-M1, M2, M3 series).

Shanghai AI Lab (InternLM)

Shanghai AI Laboratory is a national-level research institution.

InternLM3 (January 2025)

Attribute	Details
Creator	Shanghai AI Lab
Release	January 2025
License	Open-source
Size	8B (InternLM3-8B-Instruct)
Training Data	4T tokens
Key Innovation	Integrates conversational capabilities with deep thinking in single model
Efficiency	Cuts training costs by 75%+ vs comparable models

InternLM3 achieves state-of-the-art performance in knowledge understanding, reading comprehension, mathematics, and coding at its scale. However, InternLM does not have a dedicated “Code” variant – coding is integrated into the general-purpose model.

2. IDE & Coding Tools

Tongyi Lingma (Alibaba)

Attribute	Details
Creator	Alibaba Cloud
Launched	October 2023
Model Backend	Qwen series
IDE Support	VS Code, JetBrains (IntelliJ, PyCharm, GoLand, WebStorm, etc.)
Pricing	Free tier available; enterprise plans
Downloads	2M+ shortly after launch
Market Share	12.9% in China (vs GitHub Copilot’s 64.5%)
Employee ID	AI001 at Alibaba – the company’s first “AI employee”

Key Features: - Line-level and method-level code generation - Natural language to code translation - Multi-file simultaneous modification - Unit test generation, code comments, code explanation - Cross-file context awareness - Intent understanding and reflective iteration - AI-assisted code search across documentation

Alibaba projects Tongyi Lingma will write at least 20% of the company’s code going forward.

Qwen Code CLI (Alibaba)

Attribute	Details
Creator	Alibaba / Qwen Team
Launched	2025, v0.5.0 in late 2025
License	Open-source
Model Backend	Qwen3-Coder (model-agnostic, supports others)
Interface	Terminal CLI + VS Code plugin

Attribute	Details
SWE-bench	69.6%
GitHub	github.com/QwenLM/qwen-code

A direct competitor to Claude Code and Gemini CLI (it's a Gemini CLI fork). Key v0.5.0 features include: - 4 concurrent instances in a single terminal - VS Code plugin with sidebar integration - TypeScript SDK for programmatic integration - Workflow automation (PR handling, rebases, formatting) - Enhanced parser optimized for Qwen-Coder structured outputs

Kimi Code CLI (Moonshot)

Attribute	Details
Creator	Moonshot AI
Launched	January 2026 (with K2.5)
License	Apache 2.0
Model Backend	Kimi K2.5
Interface	Terminal CLI
GitHub Stars	6,400+ (as of Feb 2026)
IDE Integration	VS Code, Cursor, Zed
MCP Support	Yes

An autonomous coding agent that can analyze repos, propose plans, execute shell commands, and iterate. Supports file read/edit, command execution, code search, web content fetching, and multi-step action planning.

CodeGeeX (Zhipu AI)

Attribute	Details
Creator	Zhipu AI / Tsinghua KEG Lab
Model	CodeGeeX4-ALL-9B (latest, based on GLM-4-9B)
License	Free for individual use
IDE Support	VS Code, JetBrains (all IDEs)
HumanEval	82.3%
MBPP	75.7%
Context Window	128K tokens

Key Features: - Code completion, explanation, fixes, translation - Code interpreter, web search, function calling - Repository-level Q&A - Long-text context memory, cross-file analysis - Local mode support - NL2SQL capabilities - Automatic README generation

CodeGeeX is completely free for individuals – no token limits, no premium tier required.

Trae IDE (ByteDance)

Attribute	Details
Creator	ByteDance
Launched	January 20, 2025
Formerly	MarsCode

Attribute	Details
Base	VS Code fork
Model Backend	Doubao (ByteDance) + Claude 4, GPT-4o, DeepSeek R1
Pricing	Free tier: 5,000 auto-completions/month + premium models; Pro: \$10/month
Interface	Full IDE (desktop app)

Key Features: - **Builder Mode:** Generates full-stack apps from natural language prompts - **Chat Mode:** Step-by-step code analysis, debugging, suggestions - Multimodal input (images for UI generation) - Access to premium models (Claude 4, GPT-4o, DeepSeek R1) - VS Code extension compatibility - Strong Chinese language support (native Mandarin interface)

Privacy Concerns: Personal data retained for 5 years after account closure. Extensive telemetry collection with no opt-out. ByteDance's Unit221B blog exposed the data collection system in detail.

MarsCode Agent (the research arm) fixed 39.33% of real software bugs on SWE-bench Lite.

Baidu Comate

Attribute	Details
Creator	Baidu
Model Backend	ERNIE (Wenxin) large models
IDE Support	VS Code, JetBrains (all IDEs), Xcode, Visual Studio
Adoption	10M+ developers; 43-45% of new code at Baidu is AI-generated
Pricing	Free tier + enterprise plans
Enterprise Partners	~200 organizations testing Comate X

Comate 3.5S Features: - 5 specialized agents: code Q&A, coding, unit testing, debugging, security - **Zulu Agent:** Autonomous task decomposition and decision-making - Design-to-code conversion (F2C) - Multi-agent collaboration (agents autonomously coordinate) - MCP integration for team knowledge sharing

Comate AI IDE is the industry's first multimodal, multi-agent collaborative AI IDE, putting Baidu in direct competition with GitHub Copilot and Cognition's Devin.

Tencent CodeBuddy

Attribute	Details
Creator	Tencent Cloud
Model Backend	Hunyuan foundation model (Yuanbao Code)
Launched	Mid-2025
IDE Support	Standalone desktop app, VS Code, JetBrains
Adoption	85% of Tencent engineers use daily; 50%+ of Tencent programmers
Productivity Gain	40% reported improvement
Interface	Mandarin and English
Availability	International versions available

Features conversation-to-code workflows, allowing non-technical users to build applications via natural language. Tencent's overseas client base doubled year-over-year.

Fitten Code

Attribute	Details
Creator	Fitten Tech (Tsinghua PhD team)
Pricing	Free
IDE Support	VS Code, JetBrains, Vim, Neovim
Languages	80+ programming languages
GitHub	fittencode.nvim (Neovim), fittencode.vim (Vim)

Features real-time code completion, generation, comment generation, editing, explanation, test generation, and error detection. Represents the indie/academic tier of China’s AI coding ecosystem – built by researchers, offered free.

DevChat (Merico)

Attribute	Details
Creator	Merico
IDE Support	VS Code, JetBrains
Interface	Chat panel within IDE
Backend	OpenAI-compatible (routes to various LLMs)

Open-source tool for prompt-driven code and documentation generation. Supports smart completion, error correction, and code specification checking. Less actively developed in 2025-2026 compared to competitors.

3. Benchmarks Comparison Table

Coding Benchmarks: Chinese vs. Western Models (March 2026)

Model	Origin	HumanEval	SWE-bench Verified	LiveCodeBench	AIME 2025	Architecture
MiniMax M2.5	China	89.6%	80.2%	–	–	230B MoE (10B active)
Claude Opus 4.6	US	~92%	80.8%	–	–	Dense (proprietary)
Kimi K2.5	China	~99%	76.8%	85.0%	96.1%	1T MoE (32B active)
GLM-4.7	China	94.2%	73.8%	84.9%	95.7%	355B MoE (32B active)
Qwen3-Coder-Next	China	–	70.6%	–	–	80B MoE (3B active)
Gemini 3.1 Pro	US	93.0%	78.0%	81.3%	–	MoE (proprietary)
DeepSeek V3.2-Speciale	China	–	–	–	96.0%	671B MoE (37B active)

Model	Origin	HumanEval	SWE-bench Verified	LiveCodeBench	AIME 2025	Architecture
Qwen2.5-Coder-32B	China	92.7%	–	51.2%	–	32B Dense
GPT-5	US	–	–	–	94.6%	Proprietary
Ling-1T	China	–	–	Leading (1T class)	70.4%	1T MoE (50B active)
Yi-Coder-9B	China	85.4%	–	23.4%	–	9B Dense
ERNIE 4.5	China	>90%	–	Below leaders	–	424B MoE

Key Observations: 1. **SWE-bench Verified** (the gold standard for real-world coding): Chinese models hold 3 of the top 5 spots (MiniMax M2.5, Kimi K2.5, GLM-4.7). 2. **HumanEval** is saturating – most frontier models score >90%, making it less discriminating. 3. **Multi-SWE-Bench** (multilingual): MiniMax M2.5 leads at 51.3%, ahead of Claude Opus 4.6. 4. **Cost-adjusted performance:** Chinese models deliver comparable results at 10-30x lower cost.

4. Unique Strengths of Chinese Models

4.1 Training Efficiency & Cost

Chinese labs have pioneered techniques that fundamentally reduce training and inference costs:

- **DeepSeek’s MLA:** 93.3% reduction in KV-cache memory
- **FP8 Mixed Precision:** First validated at trillion-parameter scale by DeepSeek
- **MiniMax M2.5:** Frontier performance at 1/20th the cost of Western alternatives
- **Qwen3-Coder-Next:** Only 3B active parameters deliver 70.6% SWE-bench
- **DeepSeek V3:** Trained for \$5.6M in GPU compute vs. GPT-4’s estimated \$50-100M

This efficiency was partly born of necessity (US export controls limiting access to top-tier GPUs), but has become a genuine competitive advantage.

4.2 Open-Source Leadership

As of late 2025, four of the top five open-source models globally come from Chinese labs (MiniMax, Alibaba, DeepSeek, Z.ai). Key stats:

- Qwen has 113,000+ derivative models on HuggingFace (most of any model family)
- DeepSeek models are MIT-licensed with full commercial use rights
- Kimi K2.5 is open-weight under Modified MIT
- MiniMax M2.5, GLM-4.7, and ERNIE 4.5 are all open-source

4.3 MoE Architecture Innovation

Chinese labs have pushed MoE architecture further than Western counterparts:

Model	Total Params	Active Params	Ratio
Qwen3-Coder-Next	80B	3B	26.7:1
MiniMax M2.5	230B	10B	23:1
DeepSeek V3	671B	37B	18:1
Kimi K2.5	1T	32B	31.25:1
Ling-1T	1T	50B	20:1

Higher sparsity ratios mean more knowledge packed into fewer active compute cycles – a direct response to compute constraints.

4.4 Multilingual & Chinese-Language Superiority

Chinese models naturally excel at: - **Mandarin code comments and documentation** - **Chinese variable names and identifiers** (full Unicode support) - **Chinese-language error messages and debugging** - **Bilingual codebases** (mixed Chinese/English) - **Chinese technical documentation understanding**

This matters for the enormous domestic Chinese developer market (estimated at 7-10 million professional developers).

4.5 Agentic Coding Innovation

Several unique agentic features emerged from Chinese labs:

- **Kimi K2.5’s Agent Swarm:** 100 agents coordinating via PARL, 4.5x speedup
- **Baidu Comate’s Zulu:** Autonomous multi-agent task decomposition
- **MiniMax M2.5’s Architect Mode:** Plans features and structure before writing code
- **DeepSeek V3.2’s Thinking-in-Tool-Use:** First to integrate reasoning into tool calls

4.6 Vertical Integration

Chinese tech giants integrate coding AI across their entire stack:

Company	Model	IDE Tool	CLI Agent	Cloud Platform
Alibaba	Qwen3-Coder	Tongyi Lingma	Qwen Code	Alibaba Cloud
Baidu	ERNIE	Comate	–	Baidu Cloud
Tencent	Hunyuan	CodeBuddy	–	Tencent Cloud
ByteDance	Doubao	Trae IDE	–	Volcano Engine
Zhipu AI	GLM-4.7	CodeGeeX	–	Z.ai Platform
Moonshot	Kimi K2.5	–	Kimi Code	Kimi API

5. Access & Availability Outside China

Freely Available Globally

Model/Tool	Access Method	Notes
DeepSeek (all models)	HuggingFace, API, Ollama	MIT License, no restrictions
Qwen2.5-Coder (all sizes)	HuggingFace, ModelScope, Ollama	Apache 2.0
Qwen3-Coder / Qwen3-Coder-Next	HuggingFace, Ollama	Open-weight
Yi-Coder	HuggingFace, Ollama	Apache 2.0
Kimi K2/K2.5	HuggingFace, NVIDIA NIM, API	Modified MIT
GLM-4.7	HuggingFace, API	Open-weight
MiniMax M2/M2.5	HuggingFace, API, Ollama	Open-source
ERNIE 4.5	HuggingFace	Apache 2.0
Ling-1T	HuggingFace (inclusionAI)	Open-source
CodeGeeX plugin	VS Code / JetBrains Marketplace	Free
Trae IDE	trae.ai	Free tier globally
Fitten Code	VS Code / JetBrains / Vim Marketplace	Free

Restricted or China-Focused

Tool	Access	Notes
Tongyi Lingma	Alibaba Cloud account needed	Available internationally via Alibaba Cloud
Baidu Comate	Primarily China-focused	JetBrains plugin available globally
Tencent CodeBuddy	International version launched	Expanding globally
Qwen Code CLI	Open-source (GitHub)	Usable anywhere
Kimi Code CLI	Open-source (GitHub)	Usable anywhere

Third-Party Inference Providers

Chinese models are available through global inference providers at competitive rates:

- **SiliconFlow**: All-in-one Chinese AI inference platform (16M+ developers, 70K+ models)
- **Together AI**: Hosts Qwen3-Coder, DeepSeek, and others
- **OpenRouter**: DeepSeek V3.2, Kimi K2.5, GLM-4.7
- **NVIDIA NIM**: Kimi K2.5, MiniMax M2.5
- **Featherless**: Kimi K2.5
- **Ollama**: Most open-source Chinese models available for local inference

Integration with Western Developer Tools

Chinese models work with major Western development tools:

Tool	Chinese Model Support
Cursor	Built Composer 2 on Kimi + DeepSeek + Qwen foundations
Windsurf	Core model from Z.ai (Zhipu AI)
Continue.dev	DeepSeek, Qwen via OpenAI-compatible API
Aider	DeepSeek via API; Qwen via local/API
VS Code	CodeGeeX, Fitten Code, Tongyi Lingma, Trae plugins
JetBrains	CodeGeeX, Fitten Code, Tongyi Lingma, Comate plugins
Claude Code / Gemini CLI	Can use DeepSeek, Qwen via API configuration

6. Chinese Coding Benchmarks & Communities

Benchmarks

CRUXEval-X: An expanded multilingual code reasoning benchmark covering 19 programming languages. Developed as an extension of CRUXEval with focus on cross-lingual capabilities. Yi-Coder-9B was the first open-source model to surpass 50% on CRUXEval-O.

HumanEval Gap Closure: At the end of 2023, the performance gap on HumanEval between US and Chinese models was 31.6 percentage points. By end of 2024, this had narrowed to just 3.7 points. As of 2026, the gap has essentially vanished – Chinese models now lead on some coding benchmarks.

Benchmark Saturation: Traditional benchmarks like HumanEval are saturating (most frontier models >90%). The field is migrating to: - **SWE-bench Verified**: Real-world GitHub issue resolution - **Multi-SWE-Bench**: Multilingual real-world coding - **LiveCodeBench**: Live competitive programming - **BigCodeBench**: Only 35.5% AI success rate vs. 97% human baseline

Open-Source Communities & Platforms

ModelScope (Alibaba): - China's largest AI open-source platform - 70,000+ open-source models - 16M+ developers globally - Launched November 2022 - Functions as "Model-as-a-Service" (MaaS) - Chinese equivalent of HuggingFace

Gitee: - MIIT-backed domestic alternative to GitHub - Important for data sovereignty compliance - China's State Council proposed counting open-source contributions on Gitee toward university academic credit

HuggingFace Presence: - Chinese models are extensively published on HuggingFace - Qwen family: 113K+ derivative models (most of any family) - DeepSeek, Kimi, MiniMax, GLM, ERNIE, Ling – all have official HuggingFace organizations

7. Regulatory Considerations

Data Sovereignty

China is pushing for “algorithmic sovereignty” – complete control over computing infrastructure by 2027. Key implications for coding tools:

- Chinese companies must store code processed by AI within China
- Foreign AI coding tools used in Chinese enterprises face data residency requirements
- Domestic alternatives (Tongyi Lingma, Comate, CodeBuddy) are preferred for government and state-owned enterprise contracts

US Export Controls Impact

The Bureau of Industry and Security (BIS) maintains aggressive restrictions on GPU and accelerator exports to China:

- **Hardware restrictions:** Advanced GPUs (H100, H200) cannot be exported to China. DeepSeek V3 was trained on H800s (less restricted variant), but even these face increasing restrictions.
- **Chinese response:** “Using software to supplement hardware” – algorithmic innovation to overcome material constraints. DeepSeek’s MLA, FP8 training, and sparse attention are direct results.
- **Domestic chip development:** DeepSeek V3.1 was explicitly optimized for domestic AI accelerators (Huawei Ascend, etc.)
- **Impact on researchers:** Sharing controlled source code with foreign nationals (even within the US) can trigger licensing requirements. BIS penalties regularly exceed \$1 million.

Open-Source as Strategic Policy

China has embraced open-source AI as national strategy: - State Council policies encourage open-source contributions - Universities may count open-source work toward academic credit - Open-source models enable global adoption without direct technology transfer issues - MIT and Apache 2.0 licenses dominate Chinese model releases

Implications for Western Developers

1. **All major Chinese coding models are freely available** for download and local use – no Chinese account or residency required
 2. **API access** is available through both Chinese platforms and Western providers (OpenRouter, Together AI, SiliconFlow)
 3. **Privacy considerations:** Using Chinese-hosted APIs means data transits Chinese servers. For sensitive code, prefer local inference or Western-hosted providers
 4. **Cursor and Windsurf already use Chinese models** as backbone technology – this is not hypothetical, it’s reality
-

Summary

The Chinese AI coding ecosystem in March 2026 is not a secondary player catching up to Western leaders – it is, in several dimensions, the global leader:

- **SWE-bench Verified:** MiniMax M2.5 (80.2%) rivals Claude Opus 4.6 (80.8%). Three of the top five models are Chinese.

- **Cost efficiency:** DeepSeek is 30x cheaper than OpenAI; Qwen is 10x cheaper. MiniMax delivers frontier performance at 1/20th the cost.
- **Open-source dominance:** Four of the top five open-source models are Chinese. Qwen has more HuggingFace derivatives than any other model family.
- **Architectural innovation:** MLA, FP8 training, DeepSeek Sparse Attention, Agent Swarms, Lightning Attention – many of the most impactful recent innovations in efficient AI come from Chinese labs.
- **IDE tools:** Every major Chinese tech company has a competitive coding assistant, with several (CodeGeeX, Trae) available free globally.
- **Western adoption:** Cursor and Windsurf both build on Chinese model foundations. This is the strongest possible validation.

The competitive pressure from Chinese AI labs has materially benefited developers worldwide through lower costs, more open-source options, and faster innovation cycles. Any assessment of the AI coding tool landscape that ignores Chinese contributions is fundamentally incomplete.

Last updated: March 25, 2026